

Michael G. Noll

Applied Research. Big Data. Distributed Systems. Open Source.

- [RSS](#)

- [Blog](#)
- [Archive](#)
- [Tutorials](#)
- [Projects](#)
- [Publications](#)

Running Hadoop on Ubuntu Linux (Multi-Node Cluster)

Table of Contents

- [Tutorial approach and structure](#)
- [Prerequisites](#)
 - [Configuring single-node clusters first](#)
 - [Done? Let's continue then!](#)
- [Networking](#)
- [SSH access](#)
- [Hadoop](#)
 - [Cluster Overview \(aka the goal\)](#)
 - [Masters vs. Slaves](#)
 - [Configuration](#)
 - [conf/masters \(master only\)](#)
- [conf/slaves \(master only\)](#)
 - [conf/*-site.xml \(all machines\)](#)
 - [Additional Settings](#)
 - [Formatting the HDFS filesystem via the NameNode](#)
 - [Starting the multi-node cluster](#)
 - [HDFS daemons](#)
 - [MapReduce daemons](#)
 - [Stopping the multi-node cluster](#)
 - [MapReduce daemons](#)
 - [HDFS daemons](#)
 - [Running a MapReduce job](#)
- [Caveats](#)
 - [java.io.IOException: Incompatible namespaceIDs](#)
 - [Solution 1: Start from scratch](#)
 - [Solution 2: Manually update the namespaceID of problematic DataNodes](#)
- [Where to go from here](#)
- [Related Links](#)
- [Change Log](#)

In this tutorial I will describe the required steps for setting up a *distributed, multi-node* [Apache Hadoop](#) cluster backed by the Hadoop Distributed File System (HDFS), running on Ubuntu Linux.

Are you looking for the [single-node cluster tutorial](#)? Just [head over there](#).

Hadoop is a framework written in Java for running applications on large clusters of commodity hardware and incorporates features similar to those of the [Google File System \(GFS\)](#) and of the [MapReduce](#) computing paradigm. Hadoop's [HDFS](#) is a highly fault-tolerant distributed file system and, like Hadoop in general, designed to be deployed on low-cost hardware. It provides high throughput access to

In a previous tutorial, I described [how to setup up a Hadoop single-node cluster](#) on an Ubuntu box. The main goal of *this* tutorial is to get a more sophisticated Hadoop installation up and running, namely building a multi-node cluster using two Ubuntu boxes.

This tutorial has been tested with the following software versions:

- [Ubuntu Linux](#) 10.04 LTS (deprecated: 8.10 LTS, 8.04, 7.10, 7.04)
- [Hadoop](#) 1.0.3, released May 2012



Figure 1: Cluster of machines running Hadoop at Yahoo! (Source: Yahoo!)

Tutorial approach and structure

From two single-node clusters to a multi-node cluster - We will build a multi-node cluster using two Ubuntu boxes in this tutorial. In my humble opinion, the best way to do this for starters is to install, configure and test a "local" Hadoop setup for each of the two Ubuntu boxes, and in a second step to "merge" these two single-node clusters into one multi-node cluster in which one Ubuntu box will become the designated master (but also act as a slave with regard to data storage and processing), and the other box will become only a slave. It's much easier to track down any problems you might encounter due to the reduced complexity of doing a single-node cluster setup first on each machine.




Figure 2: Tutorial approach and structure

Let's get started!

Prerequisites

Configuring single-node clusters first

The tutorial approach outlined above means that you should read now my previous tutorial on [how to setup up a Hadoop single-node cluster](#) and follow the steps described there to build a single-node Hadoop cluster on each of the two Ubuntu boxes. It is recommended that you use the "same settings" (e.g., installation locations and paths) on both machines, or otherwise you might run into problems later when we will migrate the two machines to the final multi-node cluster setup.

Just keep in mind when [setting up the single-node clusters](#) that we will later connect and "merge" the two machines, so pick reasonable network settings etc. now for a smooth transition later.

Done? Let's continue then!

Now that you have two single-node clusters up and running, we will modify the Hadoop configuration to make one Ubuntu box the "master" (which will also act as a slave) and the other Ubuntu box a "slave".

Note: We will call the designated master machine just the "master" from now on and the slave-only machine the "slave". We will also give the two machines these respective hostnames in their networking setup, most notably in "/etc/hosts". If the hostnames of your machines are different (e.g. "node01") then you must adapt the settings in this tutorial as appropriate.

Shutdown each single-node cluster with `bin/stop-all.sh` before continuing if you haven't done so already.

Networking

This should come hardly as a surprise, but for the sake of completeness I have to point out that both machines must be able to reach each other over the network. The easiest is to put both machines in the same network with regard to hardware and software configuration, for example connect both machines via a single hub or switch and configure the network interfaces to use a common network such as 192.168.0.x/24.

To make it simple, we will assign the IP address 192.168.0.1 to the master machine and 192.168.0.2 to the slave machine. Update /etc/hosts on both machines with the following lines:

```

/etc/hosts (for master AND slave)

1 192.168.0.1   master
2 192.168.0.2   slave
  
```

SSH access

The `hduser` user on the master (aka `hduser@master`) must be able to connect a) to its own user account on the master - i.e. `ssh master` in this context and not necessarily `ssh localhost` - and b) to the `hduser` user account on the slave (aka `hduser@slave`) via a password-less SSH login. If you followed my [single-node cluster tutorial](#), you just have to add the `hduser@master`'s public SSH key (which should be in `$HOME/.ssh/id_rsa.pub`) to the `authorized_keys` file of `hduser@slave` (in this user's `$HOME/.ssh/authorized_keys`). You can do this manually or use the [following SSH command](#):

Distribute the SSH public key of `hduser@master`

```
1 hduser@master:~$ ssh-copy-id -i $HOME/.ssh/id_rsa.pub hduser@slave
```

This command will prompt you for the login password for user `hduser` on `slave`, then copy the public SSH key for you, creating the correct directory and fixing the permissions as necessary.

The final step is to test the SSH setup by connecting with user `hduser` from the master to the user account `hduser` on the slave. The step is also needed to save `slave`'s host key fingerprint to the `hduser@master`'s `known_hosts` file.

So, connecting from master to master...

```

1 hduser@master:~$ ssh master
2 The authenticity of host 'master (192.168.0.1)' can't be established.
3 RSA key fingerprint is 3b:21:b3:c0:21:5c:7c:54:2f:1e:2d:96:79:eb:7f:95.
4 Are you sure you want to continue connecting (yes/no)? yes
5 Warning: Permanently added 'master' (RSA) to the list of known hosts.
6 Linux master 2.6.20-16-386 #2 Thu Jun 7 20:16:13 UTC 2007 i686
7 ...
8 hduser@master:~$
  
```

...and from master to slave.

```

1 hduser@master:~$ ssh slave
2 The authenticity of host 'slave (192.168.0.2)' can't be established.
3 RSA key fingerprint is 74:d7:61:86:db:86:8f:31:90:9c:68:b0:13:88:52:72.
  
```

```

4 Are you sure you want to continue connecting (yes/no)? yes
5 Warning: Permanently added 'slave' (RSA) to the list of known hosts.
6 Ubuntu 10.04
7 ...
8 hduser@slave:~$

```

Hadoop

Cluster Overview (aka the goal)

The next sections will describe how to configure one Ubuntu box as a master node and the other Ubuntu box as a slave node. The master node will also act as a slave because we only have two machines available in our cluster but still want to spread data storage and processing to multiple machines.




Figure 3: How the final multi-node cluster will look like

The master node will run the “master” daemons for each layer: NameNode for the HDFS storage layer, and JobTracker for the MapReduce processing layer. Both machines will run the “slave” daemons: DataNode for the HDFS layer, and TaskTracker for MapReduce processing layer. Basically, the “master” daemons are responsible for coordination and management of the “slave” daemons while the latter will do the actual data storage and data processing work.

Masters vs. Slaves

Typically one machine in the cluster is designated as the NameNode and another machine the as JobTracker, exclusively. These are the actual “master nodes”. The rest of the machines in the cluster act as both DataNode and TaskTracker. These are the slaves or “worker nodes”.

Hadoop 1.x documentation hadoop.apache.org/common/docs/...

Configuration

conf/masters (master only)

Despite its name, the `conf/masters` file defines on which machines Hadoop will start *secondary NameNodes* in our multi-node cluster. In our case, this is just the master machine. The primary NameNode and the JobTracker will always be the machines on which you run the `bin/start-dfs.sh` and `bin/start-mapred.sh` scripts, respectively (the primary NameNode and the JobTracker will be started on the same machine if you run `bin/start-all.sh`).

Note: You can also start an Hadoop daemon manually on a machine via `bin/hadoop-daemon.sh start [namenode | secondarynamenode | datanode | jobtracker | tasktracker]`, which will not take the “`conf/masters`” and “`conf/slaves`” files into account.

Here are more details regarding the `conf/masters` file:

The secondary NameNode merges the `fsimage` and the `edits` log files periodically and keeps `edits` log size within a limit. It is usually run on a different machine than the primary NameNode since its memory requirements are on the same order as the primary NameNode. The secondary NameNode is started by “`bin/start-dfs.sh`” on the nodes specified in “`conf/masters`” file.

Hadoop HDFS user guide hadoop.apache.org/common/docs/...

Again, the machine on which `bin/start-dfs.sh` is run will become the *primary* NameNode.

On master, update `conf/masters` that it looks like this:

```

conf/masters (on master)
1 master

```

conf/slaves (master only)

The `conf/slaves` file lists the hosts, one per line, where the Hadoop slave daemons (DataNodes and TaskTrackers) will be run. We want both the master box and the slave box to act as Hadoop slaves because we want both of them to store and process data.

On master, update `conf/slaves` that it looks like this:

```

conf/slaves (on master)
1 master
2 slave

```

If you have additional slave nodes, just add them to the `conf/slaves` file, one hostname per line.

```

conf/slaves (on master)
1 master
2 slave
3 anotherslave01
4 anotherslave02

```

5 anotherslave03

Note: The `conf/slaves` file on master is used only by the scripts like `bin/start-dfs.sh` or `bin/stop-dfs.sh`. For example, if you want to add DataNodes on the fly (which is not described in this tutorial yet), you can “manually” start the DataNode daemon on a new slave machine via `bin/hadoop-daemon.sh start datanode`. Using the `conf/slaves` file on the master simply helps you to make “full” cluster restarts easier.

conf/*-site.xml (all machines)

You must change the configuration files `conf/core-site.xml`, `conf/mapred-site.xml` and `conf/hdfs-site.xml` on ALL machines as follows.

First, we have to change the [fs.default_name](#) parameter (in `conf/core-site.xml`), which specifies the [NameNode](#) (the HDFS master) host and port. In our case, this is the master machine.

`conf/core-site.xml` (ALL machines)

```
1 <property>
2   <name>fs.default.name</name>
3   <value>hdfs://master:54310</value>
4   <description>The name of the default file system. A URI whose
5   scheme and authority determine the FileSystem implementation. The
6   uri's scheme determines the config property (fs.SCHEME.impl) naming
7   the FileSystem implementation class. The uri's authority is used to
8   determine the host, port, etc. for a filesystem.</description>
9 </property>
```

Second, we have to change the [mapred.job.tracker](#) parameter (in `conf/mapred-site.xml`), which specifies the [JobTracker](#) (MapReduce master) host and port. Again, this is the master in our case.

`conf/mapred-site.xml` (ALL machines)

```
1 <property>
2   <name>mapred.job.tracker</name>
3   <value>master:54311</value>
4   <description>The host and port that the MapReduce job tracker runs
5   at. If "local", then jobs are run in-process as a single map
6   and reduce task.
7   </description>
8 </property>
```

Third, we change the [dfs.replication](#) parameter (in `conf/hdfs-site.xml`) which specifies the default block replication. It defines how many machines a single file should be replicated to before it becomes available. If you set this to a value higher than the number of available slave nodes (more precisely, the number of DataNodes), you will start seeing a lot of “(Zero targets found, forbidden1.size=1)” type errors in the log files.

The default value of `dfs.replication` is 3. However, we have only two nodes available, so we set `dfs.replication` to 2.

`conf/hdfs-site.xml` (ALL machines)

```
1 <property>
2   <name>dfs.replication</name>
3   <value>2</value>
4   <description>Default block replication.
5   The actual number of replications can be specified when the file is created.
6   The default is used if replication is not specified in create time.
7   </description>
8 </property>
```

Additional Settings

There are some other configuration options worth studying. The following information is taken from the [Hadoop API Overview](#).

In file `conf/mapred-site.xml`:

“`mapred.local.dir`”

Determines where temporary MapReduce data is written. It also may be a list of directories.

“`mapred.map.tasks`”

As a rule of thumb, use 10x the number of slaves (i.e., number of TaskTrackers).

“`mapred.reduce.tasks`”

As a rule of thumb, use `num_tasktrackers * num_reduce_slots_per_tasktracker * 0.99`. If `num_tasktrackers` is small (as in the case of this tutorial), use `(num_tasktrackers - 1) * num_reduce_slots_per_tasktracker`.

Formatting the HDFS filesystem via the NameNode

Before we start our new multi-node cluster, we must format Hadoop’s distributed filesystem (HDFS) via the NameNode. You need to do this the first time you set up an Hadoop cluster.

Warning: Do not format a running cluster because this will erase all existing data in the HDFS filesystem!

To format the filesystem (which simply initializes the directory specified by the `dfs.name.dir` variable on the NameNode), run the command

Format the cluster’s HDFS file system

```
1 hduser@master:/usr/local/hadoop$ bin/hadoop namenode -format
2 ... INFO dfs.Storage: Storage directory /app/hadoop/tmp/dfs/name has been successfully formatted.
3 hduser@master:/usr/local/hadoop$
```

Background: The HDFS name table is stored on the NameNode’s (here: `master`) local filesystem in the directory specified by `dfs.name.dir`. The name table is used by the NameNode to store tracking and coordination information for the DataNodes.

Starting the multi-node cluster

Starting the cluster is performed in two steps.

1. We begin with starting the HDFS daemons: the NameNode daemon is started on `master`, and DataNode daemons are started on all slaves (here: `master` and `slave`).
2. Then we start the MapReduce daemons: the JobTracker is started on `master`, and TaskTracker daemons are started on all slaves (here: `master` and `slave`).

HDFS daemons

Run the command `bin/start-dfs.sh` on the machine you want the (primary) NameNode to run on. This will bring up HDFS with the NameNode running on the machine you

ran the previous command on, and DataNodes on the machines listed in the `conf/slaves` file.

In our case, we will run `bin/start-dfs.sh` ON master:

Start the HDFS layer

```
1 hduser@master:/usr/local/hadoop$ bin/start-dfs.sh
2 starting namenode, logging to /usr/local/hadoop/bin/../logs/hadoop-hduser-namenode-master.out
3 slave: Ubuntu 10.04
4 slave: starting datanode, logging to /usr/local/hadoop/bin/../logs/hadoop-hduser-datanode-slave.out
5 master: starting datanode, logging to /usr/local/hadoop/bin/../logs/hadoop-hduser-datanode-master.out
6 master: starting secondarynamenode, logging to /usr/local/hadoop/bin/../logs/hadoop-hduser-secondarynamenode-master.out
7 hduser@master:/usr/local/hadoop$
```

On slave, you can examine the success or failure of this command by inspecting the log file `logs/hadoop-hduser-datanode-slave.log`.

Example output:

```
1 ... INFO org.apache.hadoop.dfs.Storage: Storage directory /app/hadoop/tmp/dfs/data is not formatted.
2 ... INFO org.apache.hadoop.dfs.Storage: Formatting ...
3 ... INFO org.apache.hadoop.dfs.DataNode: Opened server at 50010
4 ... INFO org.mortbay.util.Credential: Checking Resource aliases
5 ... INFO org.mortbay.http.HttpServer: Version Jetty/5.1.4
6 ... INFO org.mortbay.util.Container: Started org.mortbay.jetty.servlet.WebApplicationHandler@17a8a02
7 ... INFO org.mortbay.util.Container: Started WebApplicationContext[/,/]
8 ... INFO org.mortbay.util.Container: Started HttpContext[/logs,/logs]
9 ... INFO org.mortbay.util.Container: Started HttpContext[/static,/static]
10 ... INFO org.mortbay.http.SocketListener: Started SocketListener on 0.0.0.0:50075
11 ... INFO org.mortbay.util.Container: Started org.mortbay.jetty.Server@56a499
12 ... INFO org.apache.hadoop.dfs.DataNode: Starting DataNode in: FSDataSet{dirpath='/app/hadoop/tmp/dfs/data/current'}
13 ... INFO org.apache.hadoop.dfs.DataNode: using BLOCKREPORT_INTERVAL of 3538203msec
```

As you can see in slave's output above, it will automatically format its storage directory (specified by the `dfs.data.dir` parameter) if it is not formatted already. It will also create the directory if it does not exist yet.

At this point, the following Java processes should run on master...

Java processes on master after starting HDFS daemons

```
1 hduser@master:/usr/local/hadoop$ jps
2 14799 NameNode
3 15314 Jps
4 14880 DataNode
5 14977 SecondaryNameNode
6 hduser@master:/usr/local/hadoop$
```

(the process IDs don't matter of course)

...and the following on slave.

Java processes on slave after starting HDFS daemons

```
1 hduser@slave:/usr/local/hadoop$ jps
2 15183 DataNode
3 15616 Jps
4 hduser@slave:/usr/local/hadoop$
```

MapReduce daemons

Run the command `bin/start-mapred.sh` on the machine you want the JobTracker to run on. This will bring up the MapReduce cluster with the JobTracker running on the machine you ran the previous command on, and TaskTrackers on the machines listed in the `conf/slaves` file.

In our case, we will run `bin/start-mapred.sh` ON master:

Start the MapReduce layer

```
1 hduser@master:/usr/local/hadoop$ bin/start-mapred.sh
2 starting jobtracker, logging to /usr/local/hadoop/bin/../logs/hadoop-hadoop-jobtracker-master.out
3 slave: Ubuntu 10.04
4 slave: starting tasktracker, logging to /usr/local/hadoop/bin/../logs/hadoop-hduser-tasktracker-slave.out
5 master: starting tasktracker, logging to /usr/local/hadoop/bin/../logs/hadoop-hduser-tasktracker-master.out
6 hduser@master:/usr/local/hadoop$
```

On slave, you can examine the success or failure of this command by inspecting the log file `logs/hadoop-hduser-tasktracker-slave.log`. Example output:

```
1 ... INFO org.mortbay.util.Credential: Checking Resource aliases
2 ... INFO org.mortbay.http.HttpServer: Version Jetty/5.1.4
3 ... INFO org.mortbay.util.Container: Started org.mortbay.jetty.servlet.WebApplicationHandler@d19bc8
4 ... INFO org.mortbay.util.Container: Started WebApplicationContext[/,/]
5 ... INFO org.mortbay.util.Container: Started HttpContext[/logs,/logs]
6 ... INFO org.mortbay.util.Container: Started HttpContext[/static,/static]
7 ... INFO org.mortbay.http.SocketListener: Started SocketListener on 0.0.0.0:50060
8 ... INFO org.mortbay.util.Container: Started org.mortbay.jetty.Server@1e63e3d
9 ... INFO org.apache.hadoop.ipc.Server: IPC Server listener on 50050: starting
10 ... INFO org.apache.hadoop.ipc.Server: IPC Server handler 0 on 50050: starting
11 ... INFO org.apache.hadoop.mapred.TaskTracker: TaskTracker up at: 50050
12 ... INFO org.apache.hadoop.mapred.TaskTracker: Starting tracker tracker_slave:50050
13 ... INFO org.apache.hadoop.ipc.Server: IPC Server handler 1 on 50050: starting
14 ... INFO org.apache.hadoop.mapred.TaskTracker: Starting thread: Map-events fetcher for all reduce tasks on tracker_slave:50050
```

At this point, the following Java processes should run on master...

Java processes on master after starting MapReduce daemons

```
1 hduser@master:/usr/local/hadoop$ jps
2 16017 Jps
3 14799 NameNode
4 15686 TaskTracker
5 14880 DataNode
6 15596 JobTracker
7 14977 SecondaryNameNode
8 hduser@master:/usr/local/hadoop$
```

(the process IDs don't matter of course)

...and the following on slave.

Java processes on slave after starting MapReduce daemons

```
1 hduser@slave:/usr/local/hadoop$ jps
2 15183 DataNode
3 15897 TaskTracker
4 16284 Jps
5 hduser@slave:/usr/local/hadoop$
```

Stopping the multi-node cluster

Like starting the cluster, stopping it is done in two steps. The workflow however is the opposite of starting.

1. We begin with stopping the MapReduce daemons: the JobTracker is stopped on `master`, and TaskTracker daemons are stopped on all slaves (here: `master` and `slave`).
2. Then we stop the HDFS daemons: the NameNode daemon is stopped on `master`, and DataNode daemons are stopped on all slaves (here: `master` and `slave`).

MapReduce daemons

Run the command `bin/stop-mapred.sh` on the JobTracker machine. This will shut down the MapReduce cluster by stopping the JobTracker daemon running on the machine you ran the previous command on, and TaskTrackers on the machines listed in the `conf/slaves` file.

In our case, we will run `bin/stop-mapred.sh` ON `master`:

Stopping the MapReduce layer

```
1 hduser@master:/usr/local/hadoop$ bin/stop-mapred.sh
2 stopping jobtracker
3 slave: ubuntu 10.04
4 master: stopping tasktracker
5 slave: stopping tasktracker
6 hduser@master:/usr/local/hadoop$
```

Note: The output above might suggest that the JobTracker was running and stopped on “slave”, but you can be assured that the JobTracker ran on “master”.

At this point, the following Java processes should run on `master`...

Java processes on master after stopping MapReduce daemons

```
1 hduser@master:/usr/local/hadoop$ jps
2 14799 NameNode
3 18386 Jps
4 14880 DataNode
5 14977 SecondaryNameNode
6 hduser@master:/usr/local/hadoop$
```

...and the following on slave.

Java processes on slave after stopping MapReduce daemons

```
1 hduser@slave:/usr/local/hadoop$ jps
2 15183 DataNode
3 18636 Jps
4 hduser@slave:/usr/local/hadoop$
```

HDFS daemons

Run the command `bin/stop-dfs.sh` on the NameNode machine. This will shut down HDFS by stopping the NameNode daemon running on the machine you ran the previous command on, and DataNodes on the machines listed in the `conf/slaves` file.

In our case, we will run `bin/stop-dfs.sh` ON `master`:

Stopping the HDFS layer

```
1 hduser@master:/usr/local/hadoop$ bin/stop-dfs.sh
2 stopping namenode
3 slave: ubuntu 10.04
4 slave: stopping datanode
5 master: stopping datanode
6 master: stopping secondarynamenode
7 hduser@master:/usr/local/hadoop$
```

(again, the output above might suggest that the NameNode was running and stopped on `slave`, but you can be assured that the NameNode ran on `master`)

At this point, the only following Java processes should run on `master`...

Java processes on master after stopping HDFS daemons

```
1 hduser@master:/usr/local/hadoop$ jps
2 18670 Jps
3 hduser@master:/usr/local/hadoop$
```

...and the following on slave.

Java processes on slave after stopping HDFS daemons

```
1 hduser@slave:/usr/local/hadoop$ jps
2 18894 Jps
3 hduser@slave:/usr/local/hadoop$
```

Running a MapReduce job

Just follow the steps described in the section [Running a MapReduce job](#) of the [single-node cluster tutorial](#).

I recommend however that you use a larger set of input data so that Hadoop will start several Map and Reduce tasks, and in particular, on *both* `master` and `slave`. After all this installation and configuration work, we want to see the job processed by all machines in the cluster, don't we?

Here's the example input data I have used for the multi-node cluster setup described in this tutorial. I added four more Project Gutenberg etexts to the initial [three documents](#) mentioned in the single-node cluster tutorial. All etexts should be in plain text us-ascii encoding.

- [The Outline of Science, Vol. 1 \(of 4\) by J. Arthur Thomson](#)
- [The Notebooks of Leonardo Da Vinci](#)
- [Ulysses by James Joyce](#)
- [The Art of War by 6th cent. B.C. Sunzi](#)
- [The Adventures of Sherlock Holmes by Sir Arthur Conan Doyle](#)
- [The Devil's Dictionary by Ambrose Bierce](#)
- [Encyclopaedia Britannica, 11th Edition, Volume 4, Part 3](#)

Download these etexts, copy them to HDFS, run the WordCount example MapReduce job on master, and retrieve the job result from HDFS to your local filesystem.

Here's the example output on master... after executing the MapReduce job...

```
1 hduser@master:/usr/local/hadoop$ bin/hadoop jar hadoop*examples*.jar wordcount /user/hduser/gutenberg /user/hduser/gutenberg-output
2 ... INFO mapred.FileInputFormat: Total input paths to process : 7
3 ... INFO mapred.JobClient: Running job: job_0001
4 ... INFO mapred.JobClient: map 0% reduce 0%
5 ... INFO mapred.JobClient: map 28% reduce 0%
6 ... INFO mapred.JobClient: map 57% reduce 0%
7 ... INFO mapred.JobClient: map 71% reduce 0%
8 ... INFO mapred.JobClient: map 100% reduce 9%
9 ... INFO mapred.JobClient: map 100% reduce 68%
10 ... INFO mapred.JobClient: map 100% reduce 100%
11 ... INFO mapred.JobClient: Job complete: job_0001
12 ... INFO mapred.JobClient: Counters: 11
13 ... INFO mapred.JobClient: org.apache.hadoop.examples.WordCount$Counter
14 ... INFO mapred.JobClient: WORDS=1173099
15 ... INFO mapred.JobClient: VALUES=1368295
16 ... INFO mapred.JobClient: Map-Reduce Framework
17 ... INFO mapred.JobClient: Map input records=136582
18 ... INFO mapred.JobClient: Map output records=1173099
19 ... INFO mapred.JobClient: Map input bytes=6925391
20 ... INFO mapred.JobClient: Map output bytes=11403568
21 ... INFO mapred.JobClient: Combine input records=1173099
22 ... INFO mapred.JobClient: Combine output records=195196
23 ... INFO mapred.JobClient: Reduce input groups=131275
24 ... INFO mapred.JobClient: Reduce input records=195196
25 ... INFO mapred.JobClient: Reduce output records=131275
26 hduser@master:/usr/local/hadoop$
```

...and the logging output on slave for its DataNode daemon...

logs/hadoop-hduser-datanode-slave.log (on slave)

```
1 ... INFO org.apache.hadoop.dfs.DataNode: Received block blk_5693969390309798974 from /192.168.0.1
2 ... INFO org.apache.hadoop.dfs.DataNode: Received block blk_7671491277162757352 from /192.168.0.1
3 <snipp>
4 ... INFO org.apache.hadoop.dfs.DataNode: Served block blk_-7112133651100166921 to /192.168.0.2
5 ... INFO org.apache.hadoop.dfs.DataNode: Served block blk_-7545080504225510279 to /192.168.0.2
6 ... INFO org.apache.hadoop.dfs.DataNode: Served block blk_-4114464184254609514 to /192.168.0.2
7 ... INFO org.apache.hadoop.dfs.DataNode: Served block blk_-4561652742730019659 to /192.168.0.2
8 <snipp>
9 ... INFO org.apache.hadoop.dfs.DataNode: Received block blk_-2075170214887808716 from /192.168.0.2 and mirrored to /192.168.0.1:50010
10 ... INFO org.apache.hadoop.dfs.DataNode: Received block blk_1422409522782401364 from /192.168.0.2 and mirrored to /192.168.0.1:50010
11 ... INFO org.apache.hadoop.dfs.DataNode: Deleting block blk_-2942401177672711226 file /app/hadoop/tmp/dfs/data/current/blk_-2942401177672711226
12 ... INFO org.apache.hadoop.dfs.DataNode: Deleting block blk_-3019298164878756077 file /app/hadoop/tmp/dfs/data/current/blk_-3019298164878756077
```

...and on slave for its TaskTracker daemon.

logs/hadoop-hduser-tasktracker-slave.log (on slave)

```
1 ... INFO org.apache.hadoop.mapred.TaskTracker: LaunchTaskAction: task_0001_m_000000_0
2 ... INFO org.apache.hadoop.mapred.TaskTracker: LaunchTaskAction: task_0001_m_000001_0
3 ... task_0001_m_000001_0 0.08362164% hdfs://master:54310/user/hduser/gutenberg/ulyss12.txt:0+1561677
4 ... task_0001_m_000000_0 0.07951202% hdfs://master:54310/user/hduser/gutenberg/19699.txt:0+1945731
5 <snipp>
6 ... task_0001_m_000001_0 0.35611463% hdfs://master:54310/user/hduser/gutenberg/ulyss12.txt:0+1561677
7 ... Task task_0001_m_000001_0 is done.
8 ... task_0001_m_000000_0 1.0% hdfs://master:54310/user/hduser/gutenberg/19699.txt:0+1945731
9 ... LaunchTaskAction: task_0001_m_000006_0
10 ... LaunchTaskAction: task_0001_r_000000_0
11 ... task_0001_m_000000_0 1.0% hdfs://master:54310/user/hduser/gutenberg/19699.txt:0+1945731
12 ... Task task_0001_m_000000_0 is done.
13 ... task_0001_m_000006_0 0.6844295% hdfs://master:54310/user/hduser/gutenberg/132.txt:0+343695
14 ... task_0001_r_000000_0 0.095238104% reduce > copy (2 of 7 at 1.68 MB/s) >
15 ... task_0001_m_000006_0 1.0% hdfs://master:54310/user/hduser/gutenberg/132.txt:0+343695
16 ... Task task_0001_m_000006_0 is done.
17 ... task_0001_r_000000_0 0.14285716% reduce > copy (3 of 7 at 1.02 MB/s) >
18 <snipp>
19 ... task_0001_r_000000_0 0.14285716% reduce > copy (3 of 7 at 1.02 MB/s) >
20 ... task_0001_r_000000_0 0.23809525% reduce > copy (5 of 7 at 0.32 MB/s) >
21 ... task_0001_r_000000_0 0.6859089% reduce > reduce
22 ... task_0001_r_000000_0 0.7897389% reduce > reduce
23 ... task_0001_r_000000_0 0.86783284% reduce > reduce
24 ... Task task_0001_r_000000_0 is done.
25 ... Received 'KillJobAction' for job: job_0001
26 ... task_0001_r_000000_0 done; removing files.
27 ... task_0001_m_000000_0 done; removing files.
28 ... task_0001_m_000006_0 done; removing files.
29 ... task_0001_m_000001_0 done; removing files.
```

If you want to inspect the job's output data, you need to retrieve the job results from HDFS to your local file system (see instructions in the [single-node cluster tutorial](#)).

Caveats

java.io.IOException: Incompatible namespaceIDs

If you observe the error "java.io.IOException: Incompatible namespaceIDs" in the logs of a DataNode (logs/hadoop-hduser-datanode-*.log), chances are you are affected by issue [HDFS-107](#) (formerly known as [HADOOP-1212](#)).

The full error looked like this on my machines:

```

1 ... ERROR org.apache.hadoop.dfs.DataNode: java.io.IOException: Incompatible namespaceIDs in /app/hadoop/tmp/dfs/data: namenode namespaceID = 308967713; datanode namespaceID = 113030094
2   at org.apache.hadoop.dfs.DataStorage.doTransition(DataStorage.java:281)
3   at org.apache.hadoop.dfs.DataStorage.recoverTransitionRead(DataStorage.java:121)
4   at org.apache.hadoop.dfs.DataNode.startDataNode(DataNode.java:230)
5   at org.apache.hadoop.dfs.DataNode.(DataNode.java:199)
6   at org.apache.hadoop.dfs.DataNode.makeInstance(DataNode.java:1202)
7   at org.apache.hadoop.dfs.DataNode.run(DataNode.java:1146)
8   at org.apache.hadoop.dfs.DataNode.createDataNode(DataNode.java:1167)
9   at org.apache.hadoop.dfs.DataNode.main(DataNode.java:1326)

```

There are basically two solutions to fix this error as I will describe below.

Solution 1: Start from scratch

This step fixes the problem at the cost of erasing all existing data in the cluster's HDFS file system.

1. Stop the full cluster, i.e. both MapReduce and HDFS layers.
2. Delete the data directory on the problematic DataNode: the directory is specified by `dfs.data.dir` in `conf/hdfs-site.xml`; if you followed this tutorial, the relevant directory is `/app/hadoop/tmp/dfs/data`.
3. Reformat the NameNode. **WARNING: all HDFS data is lost during this process!**
4. Restart the cluster.

When deleting all the HDFS data and starting from scratch does not sound like a good idea (it might be ok during the initial setup/testing), you might give the second approach a try.

Solution 2: Manually update the namespaceID of problematic DataNodes

Big thanks to Jared Stehler for the following suggestion. This workaround is “minimally invasive” as you only have to edit a single file on the problematic DataNodes:

1. Stop the problematic DataNode(s).
2. Edit the value of `namespaceID` in `$(dfs.data.dir)/current/VERSION` to match the corresponding value of the current NameNode in `$(dfs.name.dir)/current/VERSION`.
3. Restart the fixed DataNode(s).

If you followed the instructions in my tutorials, the full paths of the relevant files are:

- NameNode: `/app/hadoop/tmp/dfs/name/current/VERSION`
- DataNode: `/app/hadoop/tmp/dfs/data/current/VERSION` (background: `dfs.data.dir` is by default set to `$(hadoop.tmp.dir)/dfs/data`, and we set `hadoop.tmp.dir` in this tutorial to `/app/hadoop/tmp`).

If you wonder how the contents of `VERSION` look like, here's one of mine:

```

contents of current/VERSION
1 namespaceID=393514426
2 storageID=D5-1706792599-10.10.10.1-50010-1204306713481
3 cTime=1215607609074
4 storageType=DATA_NODE
5 layoutVersion=-13

```

Where to go from here

If you're feeling comfortable, you can continue your Hadoop experience with my tutorial on [how to code a simple MapReduce job](#) in the Python programming language which can serve as the basis for writing your own MapReduce programs.

Related Links

From yours truly:

- [Running Hadoop On Ubuntu Linux \(Single-Node Cluster\)](#)
- [Writing An Hadoop MapReduce Program In Python](#)
- [Benchmarking and Stress Testing an Hadoop Cluster with TeraSort, TestDFSIO & Co.](#)

From other people:

- [How to debug MapReduce programs](#)
- [Hadoop API Overview](#) (Hadoop 2.x)
- [Bug HDFS-107: DataNodes should be formatted when the NameNode is formatted](#)
- [Bug MAPREDUCE-63: TaskTracker falls into an infinite loop](#) during `reduce > copy` step

Change Log

Only major changes are listed here.

- 2011-07-17: Renamed the Hadoop user from `hadoop` to `hduser` based on readers' feedback. This should make the distinction between the local Hadoop user (now `hduser`), the local Hadoop group (`hadoop`), and the Hadoop CLI tool (`hadoop`) more clear.

Tweet 4

Comments

390 Comments Michael G. Noll Login

Recommend 6 Share Sort by Best

Join the discussion...

Ivan • 3 years ago
Thanks for your tutorial. Here I got a problem.
If I put only one .txt file into dfs, the MapReduce will be completed without any problems.
But when I put more than one file, the job seems pending with below info:


```
" 12/12/08 15:51:28 INFO mapred.JobClient: map 0% reduce 0%
12/12/08 15:51:42 INFO mapred.JobClient: map 50% reduce 0%
12/12/08 15:51:43 INFO mapred.JobClient: map 100% reduce 0%
12/12/08 15:51:51 INFO mapred.JobClient: map 100% reduce 16% "
```


In the same time, output in hadoop-hduser-tasktracker-xxx.log is " ... reduce > copy (1 of 2 at 0.13MB/s)"
49 Reply Share

Sundeeep • 3 years ago
Hi Michael ,
I have followed the tutorial and created a cluster with two nodes. But when I try to run word count program. Its stopping near reduce phase. It does not proceed further. I waited for 15-20 mins and the finally quit the job. I tried many times but there no change every time its getting stuck at reduce phase. Can you suggest me what I need to do.

Regards
Sundeeep
29 Reply Share

Shuvrow • 3 years ago
Thanks Michael for your nice tutorial about hadoop. I followed your tutorial to setup hadoop in 2 pc but when i try to run word count program, Although it mapped 100%, it stop reducing after 16% or 19% completion. What can i do to complete this word count job?
Thanks
Shuvrow
16 Reply Share

Ильнур Ибатуллин Shuvrow • a year ago
I have the same problem
And I got the errors: "Task Id: attempt_201409271618_0001_m_000000_0, Status : FAILED"

Then warnings: "mapred.JobClient: Error reading task output

http://hadoop.cab407:50060/tasklog?plaintext=true&attemptid=attempt_t_201409271618_0001_m_000000_0&filter=stdout

OS: ubuntu 12.04-server LTS

hadoop 1.2.1

there were master node and 6 slave nodes(all these names were written on master node in /etc/hosts and in \$HADOOP_HOME/conf/slaves . In slave nodes there was only master's hostname.)

In master node some ports(50060, 50030, 50070, 22, 80) are opened

Where should I look for problems?
4 Reply Share

Sudheer • 3 years ago
hi

i followed the above procedure for running the multi node cluster in hadoop,
i run the jps command in master node it is showing only the following it is not showing the namenode
Jps
DataNode
SecondaryNameNode
after this again i reinstall the hadoop and now it is showing along with namenode after running the [./start-dfs.sh](#)
when i run the mapreduce [start-mapred.sh](#) then i run the jps command then it is not showing the namenode
then again i reinstall the hadoop and then at this step it is not showing the datanode

can please help me

thanks for the help....
15 Reply Share

Felipe Sudheer • 2 years ago
check your Datanode log, maybe you are facing java.io.IOException: Incompatible namespaceIDs
3 Reply Share

Aman Jot • 3 years ago
hello sir

About Me

I am a researcher and software engineer based in Switzerland, Europe. In my day job I am the developer evangelist of [Confluent](#), i.e. the US startup founded in 2014 by the creators of [Apache Kafka](#) who developed Kafka while at LinkedIn. At Confluent we are focusing on building a [stream data platform](#) to help other companies get easy access to enterprise data as real-time streams. [Read more »](#)

Contact

michael@michael-noll.com

Follow Me

[Twitter](#)
[Blog RSS](#)
[GitHub](#)
[Slideshare](#)

Recent Posts

- [Integrating Kafka and Spark Streaming: Code Examples and State of the Game](#)
- [Apache Storm 0.9 training deck and tutorial](#)
- [Apache Kafka 0.8 training deck and tutorial](#)
- [Integrating Kafka and Storm: Code Examples and State of the Game](#)
- [Wirbelsturm: 1-Click Deployments of Storm and Kafka clusters with Vagrant and Puppet](#)

Copyright © 2004-2015 [Michael G. Noll](#). All rights reserved. Views expressed here are my own. [Privacy Policy](#). Powered by [Octopress](#).